2. From Scores to Conditional Probabilities

Problem 2.1. Write $\mathbb{E}_y \left[\ell \left(y f(x) \right) \mid x \right]$ in terms of $\pi(x)$, $\ell(-f(x))$, and $\ell \left(f(x) \right)$.

Solution

$$\mathbb{E}_y \left[\ell \left(y f(x) \right) \mid x \right] = \pi(x) l(f(x)) + (1 - \pi(x)) l(-f(x))$$

Problem 2.2. Show that the Bayes prediction function $f^*(x)$ for the exponential loss function $\ell \left(y, f(x) \right) = e^{-y f(x)}$ is given by

$$f^*(x) = \frac{1}{2} \ln \left(\frac{\pi(x)}{1 - \pi(x)} \right),$$

where we've assumed $\pi(x) \in (0, 1)$. Also, show that given the Bayes prediction function f^*, we can recover the conditional probabilities by

$$\pi(x) = \frac{1}{1 + e^{-2f^*(x)}}.$$

Solution

$$f^*(x) = argmin \pi(x) e^{-f(x)} + (1 - \pi(x)) e^{f(x)}$$

Taking the derivative and setting this equal to 0 we get

$$-\pi(x) e^{-f^*(x)} + (1 - \pi(x)) e^{f^*(x)} = 0$$

$$(1 - \pi(x)) e^{f^*(x)} = \pi(x) e^{-f^*(x)}$$

$$e^{2f^*(x)} = \frac{\pi(x)}{1 - \pi(x)}$$

$$2f^*(x) = \ln \left(\frac{\pi(x)}{1 - \pi(x)} \right)$$

$$f^*(x) = \frac{1}{2} \ln \left(\frac{\pi(x)}{1 - \pi(x)} \right)$$

For the second part observe that

$$e^{2f^*(x)} = \frac{\pi(x)}{1 - \pi(x)}$$

$$e^{2f^*(x)} - \pi(x) e^{2f^*(x)} = \pi(x)$$

$$e^{2f^*(x)} = \pi(x)(1 + e^{2f^*(x)})$$

$$\frac{1}{e^{-2f^*(x)}} = \pi(x)(1 + e^{2f^*(x)})$$

$$\frac{1}{(1 + e^{-2f^*(x)})} = \pi(x)$$

Problem 2.3. Show that the Bayes prediction function $f^*(x)$ for the logistic loss function $\ell\left(y, f(x)\right) = \ln\left(1 + e^{-yf(x)}\right)$ is given by

$$f^*(x) = \ln\left(\frac{\pi(x)}{1 - \pi(x)}\right)$$

and the conditional probabilities are given by

$$\pi(x) = \frac{1}{1 + e^{-f^*(x)}}.$$

Solution

$$f^*(x) = argmin\pi(x)\ln(1 + e^{-\hat{y}}) + (1 - \pi(x))\ln(1 + e^{\hat{y}})$$

Taking the derivative and setting this equal to 0 we get

$$\frac{-\pi(x)e^{-\hat{y}}}{1 + e^{-\hat{y}}} + \frac{(1 - \pi(x))e^{\hat{y}}}{1 + e^{\hat{y}}} = 0$$

$$\frac{-\pi(x)}{1 + e^{\hat{y}}} + \frac{(1 - \pi(x))e^{\hat{y}}}{1 + e^{\hat{y}}} = 0$$

$$-\pi(x) + (1 - \pi(x))e^{\hat{y}} = 0$$

$$e^{\hat{y}} = \frac{\pi(x)}{1 - \pi(x)}$$

$$\hat{y} = f^*(x) = \ln\left(\frac{\pi(x)}{1 - \pi(x)}\right)$$

For the second part observe that

$$e^{\hat{y}} = \frac{\pi(x)}{1 - \pi(x)}$$

$$e^{\hat{y} - \pi(x)e^{\hat{y}}} = \pi(x)$$

$$e^{\hat{y}} = \pi(x)(1 + e^{\hat{y}})$$

$$\frac{e^{\hat{y}}}{1 + e^{\hat{y}}} = \pi(x)$$

$$\frac{1}{1 + e^{-\hat{y}}} = \frac{1}{1 + e^{-f^*(x)}} = \pi(x)$$

3. Logistic Regression

Problem 3.1. Show that $n\hat{R}_n(w) = \text{NLL}(w)$ for all $w \in \mathbf{R}^d$. And thus the two approaches are equivalent, in that they produce the same prediction functions.

Solution Recall that

$$n\hat{R}_n(w) = \sum_{i=1}^{n} \log\left(1 + \exp\left(-y_i w^T x_i\right)\right)$$

$$\text{NLL}(w) = \sum_{i=1}^{n} \left[-y_i' \log \phi(w^T x_i)\right] + (y_i' - 1) \log\left(1 - \phi(w^T x_i)\right)$$

Let $(x, y) \in D$ and $(x, y') \in D'$
We will consider the cases where $y = 1$ and $y = -1$ separately.

Case 1: $y = -1 \implies y' = 0$
Observe that

$$\log\left(1 + \exp\left(-y_i w^T x_i\right)\right) = \log\left(1 + \exp\left(w^T x_i\right)\right)$$

And that

$$
\begin{aligned}
\left[-y_i' \log \phi(w^T x_i)\right] + (y_i' - 1) \log\left(1 - \phi(w^T x_i)\right) &= -\log(1 - \phi(w^T x)) \\
&= -\log\left(1 - \frac{1}{1 + e^{-w^T x}}\right) \\
&= -\log\left(\frac{e^{-w^T x}}{1 + e^{-w^T x}}\right) \\
&= -\log\left(\frac{1}{1 + e^{w^T x}}\right) \\
&= \log(1 + e^{w^T x})
\end{aligned}
$$

Thus $\log\left(1 + \exp\left(-y_i w^T x_i\right)\right) = \left[-y_i' \log \phi(w^T x_i)\right] + (y_i' - 1) \log\left(1 - \phi(w^T x_i)\right)$ when $y = -1$

Case 2: $y = 1 \implies y' = 1$
Observe that

$$\log\left(1 + \exp\left(-y_i w^T x_i\right)\right) = \log\left(1 + \exp\left(-w^T x_i\right)\right)$$

And that

$$
\begin{aligned}
\left[-y_i' \log \phi(w^T x_i)\right] + (y_i' - 1) \log\left(1 - \phi(w^T x_i)\right) &= -\log(\phi(w^T x)) \\
&= \log\left(\frac{1}{1 + e^{-w^T x}}\right) \\
&= \log(1 + e^{-w^T x})
\end{aligned}
$$

Thus $\log\left(1 + \exp\left(-y_i w^T x_i\right)\right) = \left[-y_i' \log \phi(w^T x_i)\right] + (y_i' - 1) \log\left(1 - \phi(w^T x_i)\right)$ when $y = 1$

We have shown that all terms of the summations are equal. Thus we can conclude that $n\hat{R}_n(w) = \mathrm{NLL}(w)$.

Problem 3.2.1. Show that the expression for LogSumExp is valid.

Solution

$$\begin{aligned} \log(e^{x_i} + \ldots e^{x_n}) &= \log(e^{x^*}(e^{x_1 - x^*} + \ldots e^{x_n - x^*})) \\ &= \log(e^{x^*}) + \log(e^{x_1 - x^*} + \ldots e^{x_n - x^*}) \\ &= x^* + \log(e^{x_1 - x^*} + \ldots e^{x_n - x^*}) \end{aligned}$$

Problem 3.2.2. Show that $\exp\left(x_i - x^*\right) \in (0, 1]$ for any i, and thus the exp calculations will not overflow.

Solution Since $x^* = max(x_1, \ldots x_n)$, we have $x_i - x^* \leq 0$ for all i. Thus $0 < e^{x_i - x^*} \leq 1$ for all i

Problem 3.2.3. Explain why the log term in our expression $\log\left[e^{x_1-x^*} + \cdots + e^{x_n-x^*}\right]$ will never be "-inf".

Solution Note that there exists at least one x_i such that $x_i = x^*$ and that for such numbers $e^{x_i-x^*} = e^0 = 1$. Thus $e^{x_1-x^*} + \ldots e^{x_n-x^*} > 1$. Thus $\ln(e^{x_1-x^*} + \ldots e^{x_n-x^*}) > e$ and so we don't have to worry about negative infinity.

Problem 3.2.4. Show how to use the numpy function *logaddexp* to correctly compute $\log\left(1 + e^{-s}\right)$

Solution np.logaddexp(0,-s)

4. Bayesian Logistic Regression with Gaussian Priors

Problem 4.1. For the dataset $\mathcal{D}'$ described in Section 3.1, give an expression for the posterior density $p(w \mid \mathcal{D}')$ in terms of the negative log-likelihood function

Solution

$$p(w \mid D') \propto p(w)e^{-NLL(w)}$$

Problem 4.2. Suppose we take a prior on w of the form $w \sim \mathcal{N}(0, \Sigma)$. Find a covariance matrix Σ such that MAP estimate for w after observing data $\mathcal{D}'$ is the same as the minimizer of the regularized logistic regression function defined in Section 3.3 (and prove it).

Solution From section 3.3 we have

$$\hat{w} = \operatorname{argmin} \hat{R}_n(w) + \lambda \|w\|^2.$$

Now

$$
\begin{aligned}
w^* &= \operatorname{argmax} p(w \mid D') \\
&= \operatorname{argmax} p(w)e^{-NLL(w)} \\
&= \operatorname{argmin} \; -\ln\left(p(w)e^{-NLL(w)}\right) \\
&= \operatorname{argmin} \frac{1}{2}w^T\Sigma^{-1}w + NLL(w) \\
&= \operatorname{argmin} \frac{1}{2n}w^T\Sigma^{-1}w + \hat{R}_n(w) \qquad \text{since } NLL(w) = n\hat{R}_n(w)
\end{aligned}
$$

Thus

$$
\begin{aligned}
\lambda w^T w &= \frac{1}{2n}w^T\Sigma^{-1}w \\
\lambda I &= \frac{1}{2n}w^T\Sigma^{-1} \\
\Sigma &= \frac{1}{2n\lambda}I
\end{aligned}
$$

Problem 4.3. In the Bayesian approach, the prior should reflect your beliefs about the parameters before seeing the data and, in particular, should be independent on the eventual size of your dataset. Following this, you choose a prior distribution $w \sim \mathcal{N}(0, I)$. For a dataset $\mathcal{D}$ of size n, how should you choose λ in our regularized logistic regression objective function so that the minimizer is equal to the mode of the posterior distribution of w (i.e. is equal to the MAP estimator).

Solution

$$\lambda = \frac{1}{2n}$$

6. Coin Flipping Maximum Likelihood

Problem 6.1. Suppose we flip a coin and get the following sequence of heads and tails:

$$\mathcal{D} = (H, H, T)$$

Give an expression for the probability of observing $\mathcal{D}$ given that the probability of heads is θ. That is, give an expression for $p\left(\mathcal{D} \mid \theta\right)$. This is called the **likelihood of θ for the data $\mathcal{D}$**.

Solution

$$p(\mathcal{D} \mid \theta) = \theta^2(1 - \theta)$$

Problem 6.2. How many different sequences of 3 coin tosses have 2 heads and 1 tail? If we toss the coin 3 times, what is the probability of 2 heads and 1 tail? (Answer should be in terms of θ.)

Solution

$$p(2H, 1T) = 3\theta^2(1 - \theta)$$

Problem 6.3. More generally, give an expression for the likelihood $p(\mathcal{D} \mid \theta)$ for a particular sequence of flips $\mathcal{D}$ that has n_h heads and n_t tails. Make sure you have expressions that make sense even for $\theta = 0$ and $n_h = 0$, and other boundary cases. You may use the convention that $0^0 = 1$, or you can break your expression into cases if needed.

Solution

$$p(\mathcal{D} \mid \theta) = \theta^{n_h}(1 - \theta)^{n_t}$$

Problem 6.4. Show that the maximum likelihood estimate of θ given we observed a sequence with n_h heads and n_t tails is

$$\hat{\theta}_{\text{MLE}} = \frac{n_h}{n_h + n_t}.$$

You may assume that $n_h + n_t \geq 1$. (Hint: Maximizing the log-likelihood is equivalent and is often easier.)

Solution

$$\hat{\theta}_{MLE} = argmax \ \theta^{n_h}(1 - \theta)^{n_t} = argmax \ n_h \ln(\theta) + n_t \ln(1 - \theta)$$

Taking the derivative and setting it to 0 we get

$$\frac{n_h}{\hat{\theta}} - \frac{n_t}{1 - \hat{\theta}} = 0$$
$$n_h(1 - \hat{\theta}) - n_t\hat{\theta} = 0$$
$$n_h - n_h\hat{\theta} - n_t\hat{\theta} = 0$$
$$\hat{\theta} = \frac{n_h}{n_h + n_t}$$

7. Coin Flipping: Bayesian Approach with Beta Prior

Problem 7.1. Suppose that our prior distribution on θ is $\text{Beta}(h, t)$, for some $h, t > 0$. That is, $p(\theta) \propto \theta^{h-1}(1-\theta)^{t-1}$. Suppose that our sequence of flips $\mathcal{D}$ has n_h heads and n_t tails. Show that the posterior distribution for θ is $\text{Beta}(h + n_h, t + n_t)$. That is, show that

$$p(\theta \mid \mathcal{D}) \propto \theta^{h-1+n_h}(1-\theta)^{t-1+n_t}.$$

Solution

$$p(\theta) = \theta^{h-1}(1-\theta)^{t-1}$$

$$
\begin{aligned}
p(\theta \mid \mathcal{D}) &\propto (\theta)(\mathcal{D} \mid \theta) \\
&= \theta^{h-1}(1-\theta)^{t-1}\theta^{n_h}(1-\theta)^{n_t} \\
&= \theta^{h+n_h-1}(1-\theta)^{t+n_t-1}
\end{aligned}
$$

Problem 7.2. Give expressions for the MLE, the MAP, and the posterior mean estimates of θ.

Solution

$$MLE = \frac{n_h}{n_h + n_t}$$
$$MAP = \frac{h + n_h - 1}{h + t + n_h + n_t - 2}$$
$$Posterior\,Mean = \frac{h_{nh}}{h + t + n_h + n_t}$$

Problem 7.3. What happens to $\hat{\theta}_{\text{MLE}}$, $\hat{\theta}_{\text{MAP}}$, and $\hat{\theta}_{\text{POSTERIOR MEAN}}$ as the number of coin flips $n = n_h + n_t$ approaches infinity?

Solution They all approach the same value. $MLE = MAP = Posterior Mean = \frac{n_h}{n_h+n_t}$

Problem 7.4. The MAP and posterior mean estimators of θ were derived from a Bayesian perspective. Let's now evaluate them from a frequentist perspective. Suppose θ is fixed and unknown. Which of the MLE, MAP, and posterior mean estimators give **unbiased** estimates of θ, if any?

Solution The MLE is always unbiased and the MAP is unbiased only if $h = t = 1$.

$$E(MLE) = \frac{1}{n}E(n_h) = \frac{1}{n}n\theta = \theta$$

$$E(MAP) = \frac{h - 1}{n + h + t - 2} + \frac{n\theta}{n + h + t - 2} = \theta \qquad \text{when } h = t = 1$$

$$E(PosteriorMean) = \frac{h}{n + h + t} + \frac{n\theta}{n + h + t} \neq \theta \qquad \text{since } h, t > 0$$

Problem 7.5. Suppose somebody gives you a coin and asks you to give an estimate of the probability of heads, but you can only toss the coin 3 times. You have no particular reason to believe this is an unfair coin. Would you prefer the MLE or the posterior mean as a point estimate of θ? If the posterior mean, what would you use for your prior?

Solution I would prefer the posterior mean. My belief is that most coins are close to fair. With a sample size of only 3, the closest to $\theta = 0.5$ we could get is $\theta = \frac{1}{3}$ or $\theta = \frac{2}{3}$. I think that it is very unlikely that even a trick coin could be that far from fair. For a prior, I would choose a Beta distribution centered at 0.5 with low variance, something like $Beta(100, 100)$.

8. Hierarchical Bayes for Click-Trhough Rate Estimation

Problem 8.1.1. Give an expression for $p(\mathcal{D}_i \mid \theta_i)$, the likelihood of $\mathcal{D}_i$ given the probability of click θ_i, in terms of θ_i, x_i and n_i.

Solution

$$p(D_i \mid \theta_i) = \theta^{x_i}(1 - \theta)^{n_I - x_i}$$

Problem 8.1.2. We will take our prior distribution on θ_i to be Beta(a, b). The corresponding probability density function is given by

$$p(\theta_i) = \text{Beta}(\theta_i; a, b) = \frac{1}{B(a,b)} \theta_i^{a-1} (1 - \theta_i)^{b-1},$$

where $B(a, b)$ is called the Beta function. Explain (without calculation) why we must have

$$\int \theta_i^{a-1} (1 - \theta_i)^{b-1} \, d\theta_i = B(a, b).$$

Solution The area under the PDF must equal 1.

Problem 8.1.3. Give an expression for the posterior distribution $p(\theta_i \mid \mathcal{D}_i)$. In this case, include the constant of proportionality. In other words, do not use the "is proportional to" sign $\propto$ in your final expression.

Solution Note that

$$p(\theta_i \mid D_i) \propto \theta_i^{a+x_i-1}(1-\theta_i)^{n_i+b-x_i-1}$$

Now

$$\int \theta_i^{a+x_i-1}(1-\theta_i)^{n_i+b-x_i-1}d\theta_i = B(a+x_i, n_i+b-x_i)$$

Thus, in order to have a valid PDF,

$$p(\theta_i \mid D_i) = \frac{1}{B(a+x_i, n_i+b-x_i)}\theta_i^{a+x_i-1}(1-\theta_i)^{n_i+b-x_i-1}$$

Problem 8.1.4. Give a closed form expression for $p(\mathcal{D}_i)$, the marginal likelihood of $\mathcal{D}_i$, in terms of the a, b, x_i, and n_i. You may use the normalization function $B(\cdot, \cdot)$ for convenience, but you should not have any integrals in your solution.

Solution

$$
\begin{aligned}
p(D_i) &= \int p(D_i \mid \theta_i) p(\theta_i) d\theta_i \\
&= \int \frac{1}{B(a,b)} \theta_i^{a+x_i-1} (1-\theta_i)^{n_i+b-x_i-1} d\theta_i \\
&= \frac{B(a+x_i, n_i+b-x_i)}{B(a,b)}
\end{aligned}
$$

Problem 8.1.4. The maximum likelihood estimate for θ_i is x_i/n_i. Let $p_{\mathrm{MLE}}(\mathcal{D}_i)$ be the marginal likelihood of $\mathcal{D}_i$ when we use a prior on θ_i that puts all of its probability mass at x_i/n_i. Note that

$$
\begin{aligned}
p_{\mathrm{MLE}}(\mathcal{D}_i) &= p\left(\mathcal{D}_i \mid \theta_i = \frac{x_i}{n_i}\right) p\left(\theta_i = \frac{x_i}{n_i}\right) \\
&= p\left(\mathcal{D}_i \mid \theta_i = \frac{x_i}{n_i}\right).
\end{aligned}
$$

Explain why, or prove, that $p_{\mathrm{MLE}}(\mathcal{D}_i)$ is larger than $p(\mathcal{D}_i)$ for any other prior we might put on θ_i.

Solution By definition, $\theta_{MLE} = \frac{x_i}{n} = \mathrm{argmax}\, p(D_i \mid \theta)$.
Thus $p(D_i \mid \theta_{MLE}) \geq p(D_i \mid \theta)\ \forall\ \theta$
Therefore $p(\mathcal{D}_i) = \int p\left(\mathcal{D}_i \mid \theta_i\right) p(\theta_i)\, d\theta_i$ is maximized when

$$
p(\theta) = \begin{cases} 1 & \theta = \frac{x_i}{n} \\ 0 & \text{otherwise} \end{cases}
$$

Problem 8.1.5. Explain what's happening to the prior as we continue to increase the likelihood.

Solution The prior's mode approached the MLE of θ and the variance approaches 0.

logistic-regression

April 5, 2019

```
In [102]: import numpy as np
          import pandas as pd
          from sklearn.preprocessing import StandardScaler
          import matplotlib.pyplot as plt
          from sklearn.linear_model import LogisticRegression
          from sklearn.calibration import calibration_curve

In [103]: from logreg_skeleton import fit_logistic_reg, f_objective
```

3.3 Regularized Logistic Regression

3.3.1

Prove that the objective function is convex.

$J_{\text{logistic}}(w) = \frac{1}{n} \sum_{i=1}^{n} \log\left(1 + \exp\left(-y_i w^T x_i\right)\right) + \lambda \|w\|^2$

LogSumExp is convex, therefore $\log\left(1 + \exp\left(-y_i w^T x_i\right)\right)$ is convex.

Sum of convex functions is convex, therefore $\sum_{i=1}^{n} \log\left(1 + \exp\left(-y_i w^T x_i\right)\right)$ is convex.

Dividing a convex function by a positive constant is convex, therefore $\frac{1}{n} \sum_{i=1}^{n} \log\left(1 + \exp\left(-y_i w^T x_i\right)\right)$ is convex.

Norms are convex, therefore $\|w\|^2$ is convex.

Multiplying a convex function by a positive constant is convex, therefore $\lambda \|w\|^2$ is convex.

Thus $J_{\text{logistic}}(w) = \frac{1}{n} \sum_{i=1}^{n} \log\left(1 + \exp\left(-y_i w^T x_i\right)\right) + \lambda \|w\|^2$ is convex.

3.3.2

Complete the f_objective function

```
In [101]: def f_objective(theta, X, y, l2_param=1):
              '''
              Args:
                  theta: 1D numpy array of size num_features
                  X: 2D numpy array of size (num_instances, num_features)
                  y: 1D numpy array of size num_instances
                  l2_param: regularization parameter

              Returns:
                  objective: scalar value of objective function
              '''
              n = len(y)
              pred = X@theta
              margin = -y*pred
              log_loss = np.logaddexp(0, margin).sum()
              reg = l2_param*(theta@theta)
              return log_loss/n + reg
```

3.3.3

Complete the fit_logistic_reg function and use it to train a model on the procided data.

```
In [19]: def fit_logistic_reg(X, y, objective_function, l2_param=1):
             '''
             Args:
                 X: 2D numpy array of size (num_instances, num_features)
                 y: 1D numpy array of size num_instances
                 objective_function: function returning the value of the objective
                 l2_param: regularization parameter

             Returns:
                 optimal_theta: 1D numpy array of size num_features
             '''
             objective_function = partial(objective_function, X=X, y=y, l2_param=l2_param)

             n_features = X.shape[1]
             theta_0 = np.zeros(n_features)
             theta = minimize(objective_function, theta_0).x
             return theta
```

```
In [104]: # Load the data
          x_train = np.loadtxt('X_train.txt', delimiter=',')
          x_val = np.loadtxt('X_val.txt', delimiter=',')
          y_train = np.loadtxt('y_train.txt', delimiter=',')
          y_val = np.loadtxt('y_val.txt', delimiter=',')
```

```
In [105]: # Standardize data
          ss = StandardScaler()
          x_train = ss.fit_transform(x_train)
          x_val = ss.transform(x_val)
          y_train[y_train==0] = -1
          y_val[y_val==0] = -1
```

```
In [106]: # Add bias term
          x_train = np.append(10*np.ones((len(x_train),1)), x_train, axis=1)
          x_val = np.append(10*np.ones((len(x_val),1)), x_val, axis=1)
```

```
In [107]: # Train model
          theta = fit_logistic_reg(x_train, y_train, f_objective)
```

```
In [108]: theta
```

```
Out[108]: array([ 0.00236204,  0.00095657, -0.00030132,  0.00302058,  0.10533832,
                -0.00358714, -0.00135921, -0.00385466, -0.00079028, -0.0011443 ,
                -0.07179551,  0.00655072, -0.004512  ,  0.01125831, -0.003866  ,
                -0.00271356,  0.00150264, -0.00278385, -0.00919238, -0.00682348,
                -0.01027393])
```

3.3.4

Find the l2 regulrization term that minimizes the log-likelihood on the validation set. Plot the log-likelihood for different values.

```
In [113]: def NLL(x, y, theta):
              margin = -y * (x@theta)
              return np.logaddexp(0, margin).sum()

In [122]: nlls = []
          l2s = np.logspace(-5, -1, 40)
          for l2 in l2s:
              theta = fit_logistic_reg(x_train, y_train, f_objective, l2)
              nlls.append(NLL(x_val, y_val, theta))

In [123]: plt.plot(l2s, nlls)
          plt.xscale('log')
          plt.xlabel('l2 Regularization')
          plt.ylabel('NLL')

Out[123]: Text(0, 0.5, 'NLL')
```


Based on this, I'll choose 0.011 as my l2 reg

3.3.5

Calibration plot

```
In [124]: def sigmoid(x):
              return 1/(1+np.exp(-x))

In [125]: theta = fit_logistic_reg(x_train, y_train, f_objective, 0.011)
          y_pred = sigmoid(x_val@theta)

In [146]: fraction_of_positives, mean_predicted_value = calibration_curve(y_val, y_pred, n_bins=

In [147]: plt.plot(mean_predicted_value, fraction_of_positives)
          plt.xlabel('mean_predicted_value')
          plt.ylabel('fraction_of_positives')

Out[147]: Text(0, 0.5, 'fraction_of_positives')
```


It appears that the model unerconfident as hown by the sigmoid shape of the calibration curve. Sigmoid calibration may help correct and calibrate the predictions.

bayesian_regression_support

April 5, 2019

0.1 Recreating figure 3.7 from Bishop's "Pattern Recognition and Machine Learning."

This notebook provides scaffolding for your exploration Bayesian Linear Gaussian Regression, as described in Lecture. In particular, through this notebook you will reproduce several variants of figure 3.7 from Bishop's book.

0.2 Instructions:

0.2.1 5.1-3:

Implement the functions in `problem` -- completed implementations of these functions are needed to generate the plots.

```
In [3]: from support_code import *
        from problem import *
```

0.3 Instructions (continued):

0.3.1 5.4:

If your implementations are correct, then the next few code blocks in this notebook will generate the required variants of Bishop's figure. These are the same figures that you would obtain if you ran `python problem.py` from the command line -- this notebook is just provided as additional support.

```
In [4]: # Generate our simulated dataset
        # Note we are using sigma == 0.2

        np.random.seed(46134)
        actual_weights = np.matrix([[0.3], [0.5]])
        data_size = 40
        noise = {"mean":0, "var":0.2 ** 2}
        likelihood_var = noise["var"]
        xtrain, ytrain = generate_data(data_size,
                                       noise,
                                       actual_weights)
```

Next, we generate the plots using 3 different prior covariance matrix. In the main call to `problem.py`, this is done in a loop -- here we wrap the loop body in a short helper function.

```
In [9]: def make_plot_given_sigma(sigma_squared):
            prior = {"mean":np.matrix([[0], [0]]),
                     "var":matlib.eye(2) * sigma_squared}

            make_plots(actual_weights,
                       xtrain,
                       ytrain,
                       likelihood_var,
                       prior,
                       likelihood_func,
                       get_posterior_params,
                       get_predictive_params,
                       sigma_squared)

In [8]: sigmas = [1/2, 1/(2**5), 1/(2**10)]
```
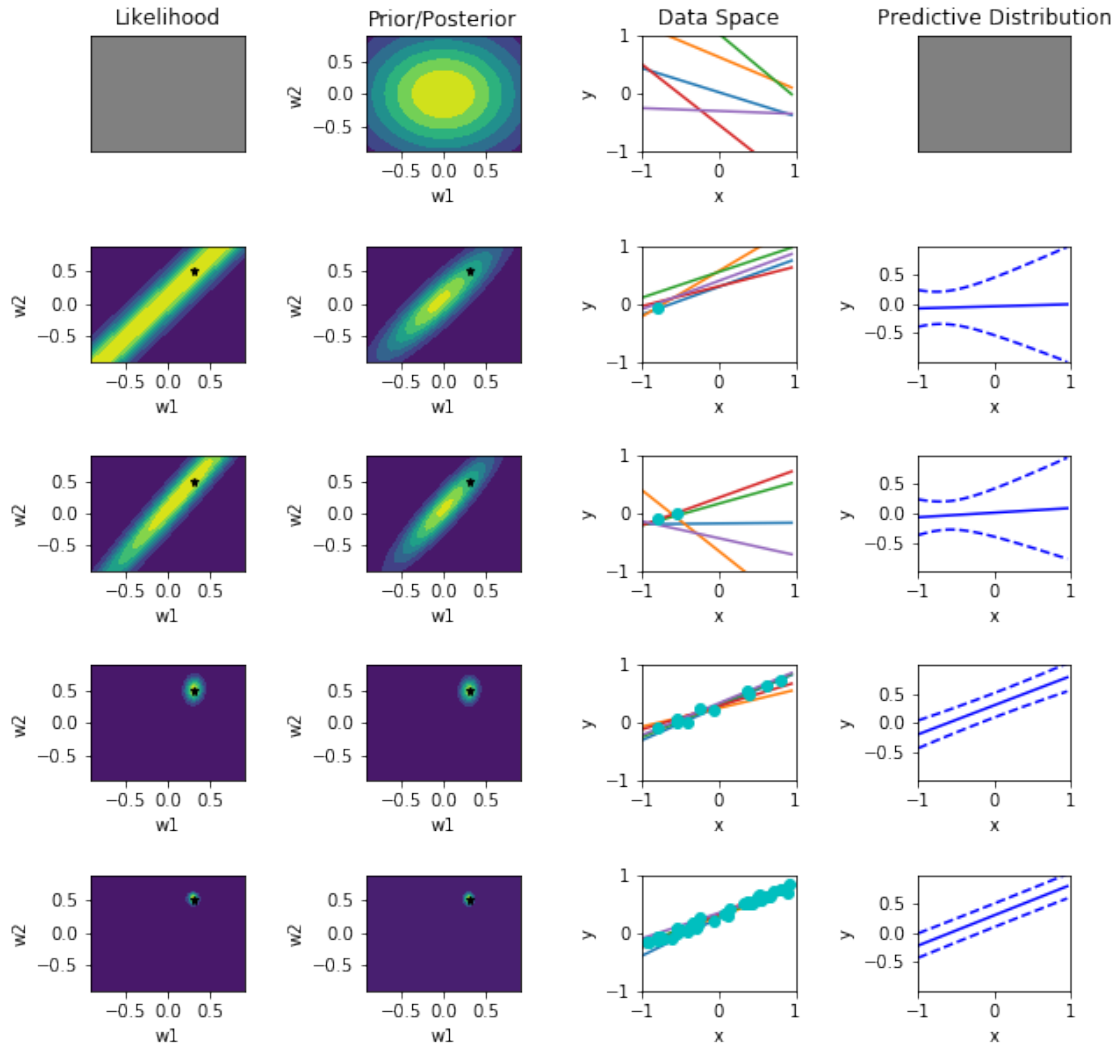
First covariance matrix:
$$\Sigma_0 = \frac{1}{2}I, \qquad I \in \mathbb{R}^{2 \times 2}$$

```
In [7]: try:
            make_plot_given_sigma(sigmas[0])
        except NameError:
            print('If not yet implemented, implement functions in problem.py.')
            print('If you have implemented, remove this try/except.')
```

Second covariance matrix:

$$\Sigma_0 = \frac{1}{2^5}I, \qquad I \in \mathbb{R}^{2\times 2}$$

```
In [8]: try:
            make_plot_given_sigma(sigmas[1])
        except NameError:
            print('If not yet implemented, implement functions in problem.py.')
            print('If you have implemented, remove this try/except.')
```

Third covariance matrix:

$$\Sigma_0 = \frac{1}{2^{10}} I, \qquad I \in \mathbb{R}^{2\times 2}$$

```
In [9]: try:
            make_plot_given_sigma(sigmas[2])
        except NameError:
            print('If not yet implemented, implement functions in problem.py.')
            print('If you have implemented, remove this try/except.')
```

0.4   Instructions (continued):

0.4.1   5.5:

The likelihood function is unaffected by the strength of the prior. As more samples are drawn, the likelihood function becomes more centered around the true parameters and its variance decreases.

The posterior distribution is strongly affected by the strength of the prior. In the final plot, the prior is so strong, that even after many observations, our parameter distribution has not come close to the true values.

The predictive distribution is also strongly affected by the strength of the prior. With a weak prior, there is high variance around unseen values. With a strong prior, although the variance is reduced, the predictive distribution is not able to converge to the true values.

0.5   Instructions (continued):

0.5.1   5.6:

For question (6), find the MAP solution for the first prior covariance $\left(\frac{1}{2}I\right)$ by completing the implementation below. In addition, be sure to justify the value for the regularization coefficient (in sklearn named alpha) in your written work.

```python
In [1]: from sklearn.linear_model import Ridge

In [15]: n = len(ytrain)
         alpha = 1/(2*n*sigmas[0])
         #alpha = 9999 # Change to the correct value
         ridge = Ridge(alpha=alpha,
                       fit_intercept=False,
                       solver='cholesky')

         ridge.fit(xtrain, ytrain)

Out[15]: Ridge(alpha=0.025, copy_X=True, fit_intercept=False, max_iter=None,
               normalize=False, random_state=None, solver='cholesky', tol=0.001)
```

   If alpha is set correctly, ridge.coef_ will equal the prior mean/MAP estimate returned by the next two cells.

```python
In [16]: ridge.coef_

Out[16]: array([[0.30085783, 0.52614072]])

In [14]: prior = {"mean":np.matrix([[0], [0]]),
                  "var":matlib.eye(2) * sigmas[0]}


         post = get_posterior_params(xtrain, ytrain, prior,
                                     likelihood_var = 0.2**2)
         post[0].ravel()

Out[14]: matrix([[0.30052135, 0.52406189]])
```