

1. Reformulations of Multiclass Hinge Loss

Problem 1.2.1. Show that if $\Delta(y, y) = 0$ for all $y \in \mathcal{Y}$, then $\ell_2(h, (x_i, y_i)) = \ell_1(h, (x_i, y_i))$.

Solution Assume $\Delta(y, y) = 0$ for all $y \in \mathcal{Y}$, then

$$\begin{aligned}
 \ell_2(h, (x_i, y_i)) &= \max_{y \in \mathcal{Y}} [\Delta(y_i, y) + h(x_i, y) - h(x_i, y_i)] \\
 &= \max \left[\Delta(y_i, y_i) + h(x_i, y_i) - h(x_i, y_i), \max_{y \neq y_i} [\Delta(y_i, y) + h(x_i, y) - h(x_i, y_i)] \right] \\
 &= \max \left[0, \max_{y \neq y_i} [\Delta(y_i, y) + h(x_i, y) - h(x_i, y_i)] \right] \\
 &= \max_{y \neq y_i} [\max[0, \Delta(y_i, y) + h(x_i, y) - h(x_i, y_i)]] \\
 &= \ell_1(h, (x_i, y_i))
 \end{aligned}$$

Problem 1.2.2.a. Show that under the conditions above, $\ell_1(h, (x_i, y_i)) = \ell_2(h, (x_i, y_i)) = 0$.

Solution Since $\Delta(y, y) = 0$, $\ell_1 = \ell_2$.

Also, since $m_{i,y}(h) = h(x_i, y_i) - h(x_i, y) \geq \Delta(y_i, y)$

$$\Delta(y_i, y) - m_{i,y}(h) = \Delta(y_i, y) + h(x_i, y) - h(x_i, y_i) \leq 0 \quad \forall y \neq y_i$$

Then it is clear that $\ell_1(h, (x_i, y_i)) = 0 = \ell_2(h, (x_i, y_i))$

Problem 1.2.2.b. Show that under the conditions above, we make the correct prediction on x_i . That is, $f(x_i) = \arg \max_{y \in \mathcal{Y}} h(x_i, y) = y_i$.

Solution Assume $f(x_i) = \arg \max_{y \in \mathcal{Y}} h(x_i, y) \neq y_i$.

Then $\exists y'$ such that $h(x_i, y') > h(x_i, y_i)$.

Then $h(x_i, y_i) - h(x_i, y) < 0$. But this contradicts the fact that

$$h(x_i, y_i) - h(x_i, y') \geq \Delta(y_i, y) > 0$$

Thus we conclude that $f(x_i) = \arg \max_{y \in \mathcal{Y}} h(x_i, y) = y_i$

2. SGD for Multiclass Linear SV

Problem 2.2. Since $J(w)$ is convex, it has a subgradient at every point. Give an expression for a subgradient of $J(w)$. You may use any standard results about subgradients, including the result from an earlier homework about subgradients of the pointwise maxima of functions.

Solution

$$\Delta J(w) = 2\lambda w + \frac{1}{n} \sum_{i=1}^n [\Psi(x_i, \hat{y}_i) - \Psi(x_i, y_i)]$$

Problem 2.3. Give an expression for the stochastic subgradient based on the point (x_i, y_i) .

Solution

$$\Delta J(w) = 2\lambda w + \Psi(x_i, \hat{y}_i) - \Psi(x_i, y_i)$$

Problem 2.4.

Give an expression for a minibatch subgradient, based on the points $(x_i, y_i), \dots, (x_{i+m-1}, y_{i+m-1})$.

Solution

$$\Delta J(w) = 2\lambda w + \frac{1}{m} \sum_{i=1}^m [\Psi(x_i, \hat{y}_i) - \Psi(x_i, y_i)]$$

3. Hinge Loss is a Special Case of Generalized Hinge Loss

Problem 3. Let $\mathcal{Y} = \{-1, 1\}$. Let $\Delta(y, \hat{y}) = 1(y \neq \hat{y})$. If $g(x)$ is the score function in our binary classification setting, then define our compatibility function as

$$\begin{aligned} h(x, 1) &= g(x)/2 \\ h(x, -1) &= -g(x)/2. \end{aligned}$$

Show that for this choice of h , the multiclass hinge loss reduces to hinge loss:

$$\ell(h, (x, y)) = \max_{y' \in \mathcal{Y}} [\Delta(y, y') + h(x, y') - h(x, y)] = \max\{0, 1 - yg(x)\}$$

Solution Note that

$$\ell(h, (x, y)) = \max[\Delta(-1, y') + h(x, y') - h(x, -1), \Delta(1, y') + h(x, y') - h(x, 1)]$$

Either $y = y'$ or $y \neq y'$.

If $y = y'$, then $\ell(h, (x, y)) = 0$.

Otherwise

$$\begin{aligned} \ell(h, (x, y)) &= \Delta(y, y') + h(x, y') - h(x, y) \\ &= \begin{cases} 1 + g(x) & \text{if } y = -1 \\ 1 - g(x) & \text{if } y = 1 \end{cases} \\ &= 1 - yg(x) \end{aligned}$$

Thus $\ell(h, (x, y)) = \max\{0, 1 - yg(x)\}$

Gradient Boosting Machines

Problem 7.1. Consider the regression framework, where $\mathcal{Y} = \mathbf{R}$. Suppose our loss function is given by

$$\ell(\hat{y}, y) = \frac{1}{2} (\hat{y} - y)^2,$$

and at the beginning of the m 'th round of gradient boosting, we have the function $f_{m-1}(x)$. Show that the h_m chosen as the next basis function is given by

$$h_m = \arg \min_{h \in \mathcal{F}} \sum_{i=1}^n [(y_i - f_{m-1}(x_i)) - h(x_i)]^2.$$

In other words, at each stage we find the base prediction function $h_m \in \mathcal{F}$ that is the best fit to the residuals from the previous stage.

Solution Note that

$$\begin{aligned} (g_m)_i &= \frac{\partial}{\partial f(x_i)} \sum_{j=1}^n \ell(y_j, f(x_j)) \\ &= \frac{\partial}{\partial f(x_i)} \frac{1}{2} (y_i - f(x_i))^2 \\ &= f(x_i) - y_i \end{aligned}$$

Thus $h_m = \arg \min_{h \in \mathcal{F}} \sum_{i=1}^n [(y_i - f_{m-1}(x_i)) - h(x_i)]^2$.

Problem 7.2. Now let's consider the classification framework, where $\mathcal{Y} = \{-1, 1\}$. In lecture, we noted that AdaBoost corresponds to forward stagewise additive modeling with the exponential loss, and that the exponential loss is not very robust to outliers (i.e. outliers can have a large effect on the final prediction function). Instead, let's consider the logistic loss

$$\ell(m) = \ln(1 + e^{-m}),$$

where $m = yf(x)$ is the margin. Similar to what we did in the ℓ_2 -Boosting question, write an expression for h_m as an argmin over \mathcal{F} .

Solution Note that $\ell(y, f(x)) = \ln(1 + e^{-yf(x)})$.

Then

$$\begin{aligned} (g_m)_i &= \frac{\partial}{\partial f(x_i)} \sum_{j=1}^n \ln(1 + e^{-y_j f(x_j)}) \\ &= \frac{\partial}{\partial f(x_i)} \ln(1 + e^{-y_i f(x_i)}) \\ &= \frac{-y_i e^{-y_i f(x_i)}}{1 + e^{-y_i f(x_i)}} \end{aligned}$$

Thus $h_m = \arg \min_{h \in \mathcal{F}} \sum_{i=1}^n \left[\frac{-y_i e^{-y_i f(x_i)}}{1 + e^{-y_i f(x_i)}} - h(x_i) \right]^2$

CART-GBM-skeleton-code

April 29, 2019

```
In [23]: import matplotlib.pyplot as plt
         from itertools import product
         import numpy as np
         from collections import Counter
         from sklearn.base import BaseEstimator, RegressorMixin, ClassifierMixin
         from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, export_graphviz
         import graphviz
         import pdb
         from IPython.display import Image

         %matplotlib inline
```

1 Load Data

```
In [24]: data_train = np.loadtxt('svm-train.txt')
         data_test = np.loadtxt('svm-test.txt')
         x_train, y_train = data_train[:, 0: 2], data_train[:, 2].reshape(-1, 1)
         x_test, y_test = data_test[:, 0: 2], data_test[:, 2].reshape(-1, 1)

In [25]: # Change target to 0-1 label
         y_train_label = np.array(list(map(lambda x: 1 if x > 0 else 0, y_train))).reshape(-1, 1)
```

2 Decision Tree Class

```
In [26]: class Decision_Tree(BaseEstimator):

         def __init__(self, split_loss_function, leaf_value_estimator,
                     depth=0, min_sample=5, max_depth=10):
             '''
             Initialize the decision tree classifier

             :param split_loss_function: method for splitting node
             :param leaf_value_estimator: method for estimating leaf value
             :param depth: depth indicator, default value is 0, representing root node
             :param min_sample: an internal node can be splitted only if it contains points
             :param max_depth: restriction of tree depth.
             '''
```

```

self.split_loss_function = split_loss_function
self.leaf_value_estimator = leaf_value_estimator
self.depth = depth
self.min_sample = min_sample
self.max_depth = max_depth

def is_pure(self, y):
    return len(set(y.flatten())) <= 1

def fit(self, X, y=None):
    """
    This should fit the tree classifier by setting the values self.is_leaf,
    self.split_id (the index of the feature we want ot split on, if we're splitting
    self.split_value (the corresponding value of that feature where the split is),
    and self.value, which is the prediction value if the tree is a leaf node. If u
    splitting the node, we should also init self.left and self.right to be Decision
    objects corresponding to the left and right subtrees. These subtrees should be
    the data that fall to the left and right, respectively, of self.split_value.
    This is a recursive tree building procedure.

    :param X: a numpy array of training data, shape = (n, m)
    :param y: a numpy array of labels, shape = (n, 1)

    :return self
    """

    if self.depth >= self.max_depth or len(y) <= self.min_sample or self.is_pure(y):
        self.is_leaf = True
        self.value = self.leaf_value_estimator(y)

    else:
        self.is_leaf = False
        splits = [] # contains (feature_index, split_value, loss) tuples
        for i, x in enumerate(X.T): # iterate over columns
            losses = [] # contain (split_value, loss) pairs
            for split_val in x:
                left_y = y[x<=split_val]
                right_y = y[x>split_val]
                loss = len(left_y)*self.split_loss_function(left_y) + len(right_y)*
                losses.append([split_val, loss])
            min_loss = min(losses, key=lambda x: x[1])
            splits.append([i] + min_loss)

        #pdb.set_trace()
        min_split = min(splits, key=lambda x: x[2])
        self.split_id = min_split[0]
        self.split_value = min_split[1]

```

```

left_x = X[X[:,self.split_id] <= self.split_value]
right_x = X[X[:,self.split_id] > self.split_value]
left_y = y[X[:,self.split_id] <= self.split_value]
right_y = y[X[:,self.split_id] > self.split_value]

self.left = Decision_Tree(self.split_loss_function,
                           self.leaf_value_estimator,
                           self.depth + 1,
                           self.min_sample,
                           self.max_depth)

self.right = Decision_Tree(self.split_loss_function,
                            self.leaf_value_estimator,
                            self.depth + 1,
                            self.min_sample,
                            self.max_depth)

#pdb.set_trace()
try:
    self.left.fit(left_x, left_y)
    self.right.fit(right_x, right_y)
except:
    pdb.set_trace()

return self

def predict_instance(self, instance):
    '''
    Predict label by decision tree

    :param instance: a numpy array with new data, shape (1, m)

    :return whatever is returned by leaf_value_estimator for leaf containing instance
    '''
    if self.is_leaf:
        return self.value
    if instance[self.split_id] <= self.split_value:
        return self.left.predict_instance(instance)
    else:
        return self.right.predict_instance(instance)

```

3 Decision Tree Classifier

```
In [27]: def compute_entropy(label_array):
        '''
        Calculate the entropy of given label list

        :param label_array: a numpy array of labels shape = (n, 1)
        :return entropy: entropy value
        '''
        counter = Counter(label_array.flatten())
        entropy = 0
        # pdb.set_trace()
        for c, n in counter.items():
            p_c = float(n)/len(label_array)
            entropy -= p_c*np.log2(p_c)
        return entropy

def compute_gini(label_array):
    '''
    Calculate the gini index of label list

    :param label_array: a numpy array of labels shape = (n, 1)
    :return gini: gini index value
    '''
    gini = 0
    for c, n in Counter(label_array.flatten()).items():
        p_c = float(n)/len(label_array)
        # gini += p_c * (1-p_c)
        gini += p_c**2
    return 1 - gini

In [50]: def most_common_label(y):
        '''
        Find most common label
        '''
        label_cnt = Counter(y.reshape(len(y)))
        #pdb.set_trace()

        label = label_cnt.most_common(1)[0][0]
        #pdb.set_trace()

        return label

In [51]: class Classification_Tree(BaseEstimator, ClassifierMixin):

        loss_function_dict = {
            'entropy': compute_entropy,
```

```

        'gini': compute_gini
    }

def __init__(self, loss_function='entropy', min_sample=5, max_depth=10):
    '''
    :param loss_function(str): loss function for splitting internal node
    '''

    self.tree = Decision_Tree(self.loss_function_dict[loss_function],
                              most_common_label,
                              0, min_sample, max_depth)

def fit(self, X, y=None):
    self.tree.fit(X,y)
    return self

def predict_instance(self, instance):
    value = self.tree.predict_instance(instance)
    return value

```

4 Decision Tree Boundary

```
In [80]: # Training classifiers with different depth
clf1 = Classification_Tree(max_depth=1)
clf1.fit(x_train, y_train_label)

clf2 = Classification_Tree(max_depth=2)
clf2.fit(x_train, y_train_label)

clf3 = Classification_Tree(max_depth=3)
clf3.fit(x_train, y_train_label)

clf4 = Classification_Tree(max_depth=4)
clf4.fit(x_train, y_train_label)

clf5 = Classification_Tree(max_depth=5)
clf5.fit(x_train, y_train_label)

clf6 = Classification_Tree(max_depth=6)
clf6.fit(x_train, y_train_label)

# Plotting decision regions
x_min, x_max = x_train[:, 0].min() - 1, x_train[:, 0].max() + 1
y_min, y_max = x_train[:, 1].min() - 1, x_train[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

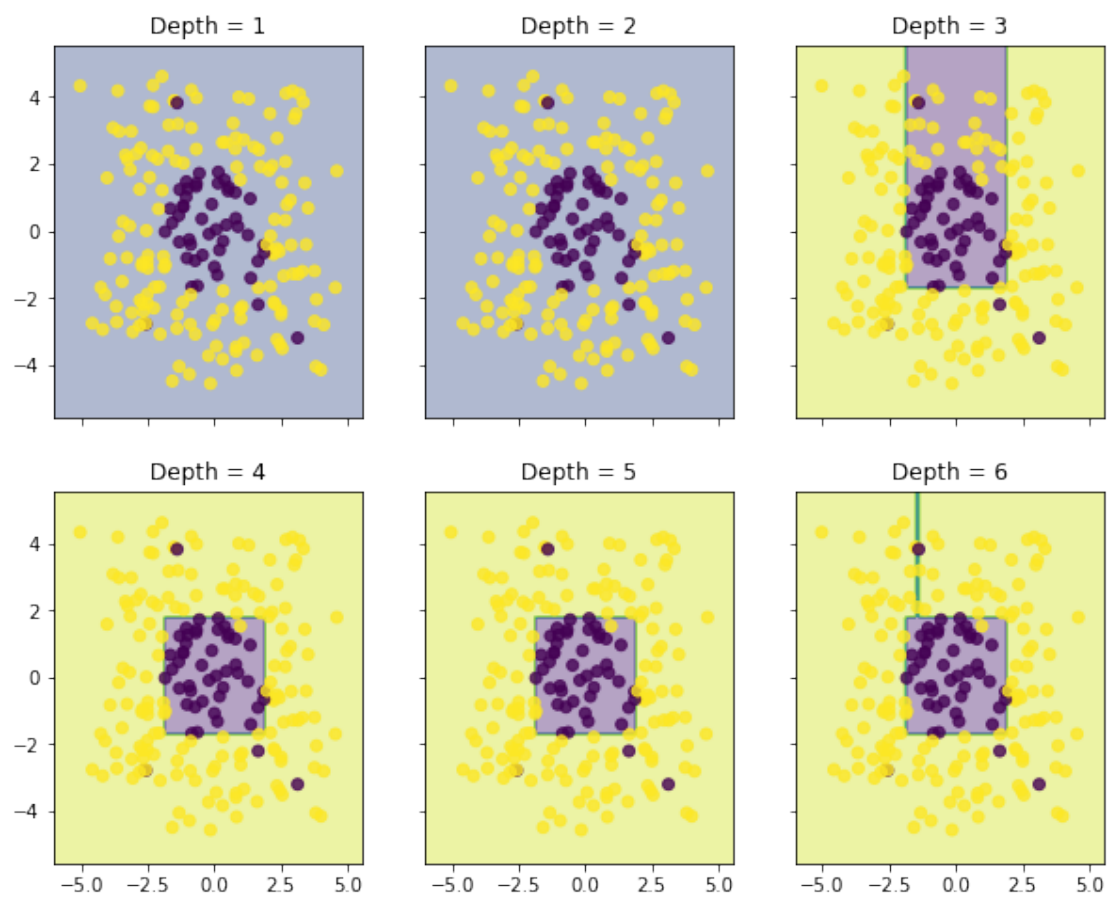
f, axarr = plt.subplots(2, 3, sharex='col', sharey='row', figsize=(10, 8))

for idx, clf, tt in zip(product([0, 1], [0, 1, 2]),
                        [clf1, clf2, clf3, clf4, clf5, clf6],
                        ['Depth = {}'.format(n) for n in range(1, 7)]):

    Z = np.array([clf.predict_instance(x) for x in np.c_[xx.ravel(), yy.ravel()]])
    Z = Z.reshape(xx.shape)
    #pdb.set_trace

    axarr[idx[0], idx[1]].contourf(xx, yy, Z, alpha=0.4)
    axarr[idx[0], idx[1]].scatter(x_train[:, 0], x_train[:, 1], c=y_train_label.flatten())
    axarr[idx[0], idx[1]].set_title(tt)

plt.show()
```



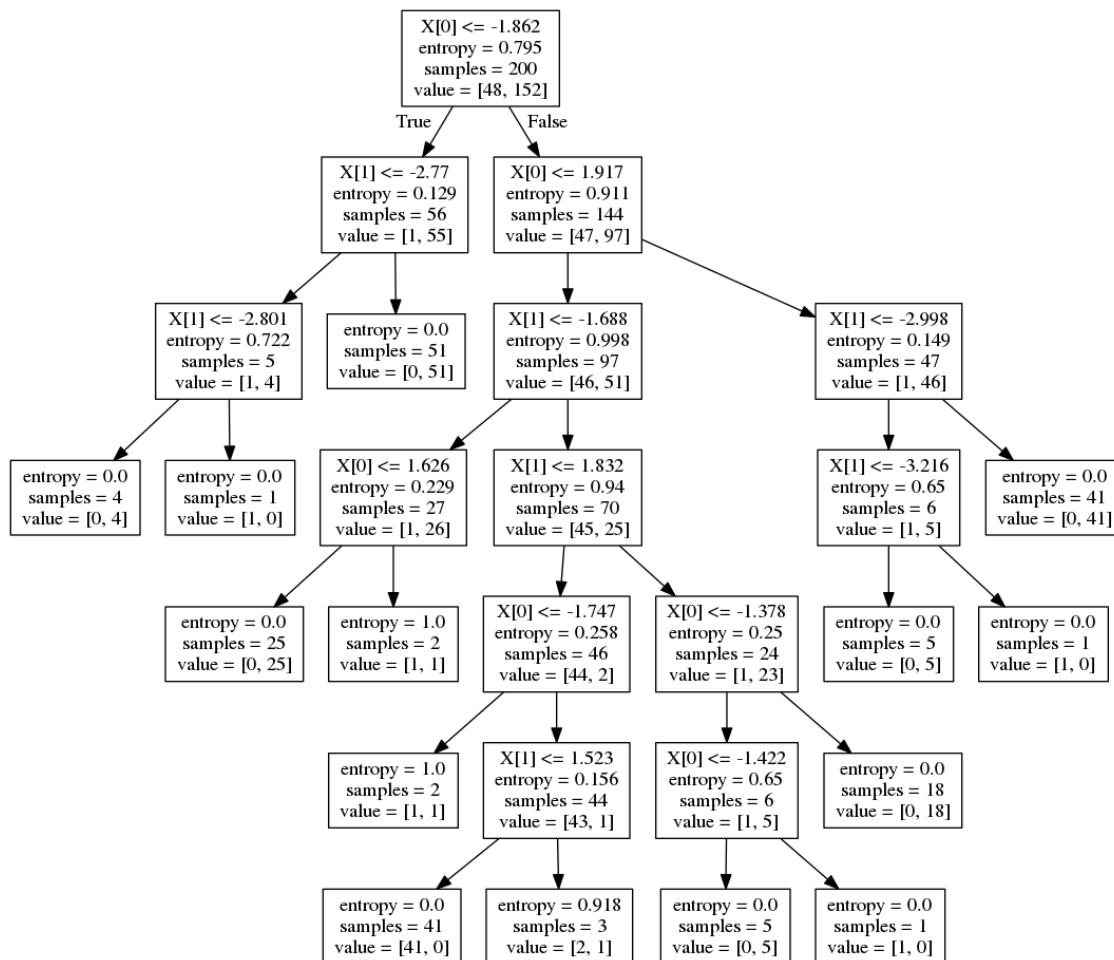
5 Compare decision tree with tree model in sklearn

```
In [75]: clf = DecisionTreeClassifier(criterion='entropy', max_depth=6, min_samples_split=5)
         clf.fit(x_train, y_train_label)
         export_graphviz(clf, out_file='tree_classifier.dot')
```

```
In [76]: # Visualize decision tree
         !dot -Tpng tree_classifier.dot -o tree_classifier.png
```

```
In [77]: Image(filename='tree_classifier.png')
```

Out[77]:



```
In [46]: clf2.tree.left.right.value
```

Out[46]: 1

6 Decision Tree Regressor

In [67]: *# Regression Tree Specific Code*

```
def mean_absolute_deviation_around_median(y):  
    '''  
        Calculate the mean absolute deviation around the median of a given target list  
  
    :param y: a numpy array of targets shape = (n, 1)  
    :return mae  
    '''  
    mean = np.mean(y)  
    n = len(y)  
    mae = np.abs(y - mean).sum()/float(n)  
    return mae
```

In [68]: `class Regression_Tree():`

```
    '''  
    :attribute loss_function_dict: dictionary containing the loss functions used for sp  
    :attribute estimator_dict: dictionary containing the estimation functions used in l  
    '''  
  
    loss_function_dict = {  
        'mse': np.var,  
        'mae': mean_absolute_deviation_around_median  
    }  
  
    estimator_dict = {  
        'mean': np.mean,  
        'median': np.median  
    }  
  
    def __init__(self, loss_function='mse', estimator='mean', min_sample=5, max_depth=1  
        '''  
        Initialize Regression_Tree  
        :param loss_function(str): loss function used for splitting internal nodes  
        :param estimator(str): value estimator of internal node  
        '''  
  
        self.tree = Decision_Tree(self.loss_function_dict[loss_function],  
                                   self.estimator_dict[estimator],  
                                   0, min_sample, max_depth)  
  
    def fit(self, X, y=None):  
        self.tree.fit(X,y)  
        return self  
  
    def predict_instance(self, instance):  
        value = self.tree.predict_instance(instance)
```

```
return value
```

7 Fit regression tree to one-dimensional regression data

```
In [69]: data_krr_train = np.loadtxt('krr-train.txt')
data_krr_test = np.loadtxt('krr-test.txt')
x_krr_train, y_krr_train = data_krr_train[:,0].reshape(-1,1), data_krr_train[:,1].reshape(-1,1)
x_krr_test, y_krr_test = data_krr_test[:,0].reshape(-1,1), data_krr_test[:,1].reshape(-1,1)

# Training regression trees with different depth
clf1 = Regression_Tree(max_depth=1, min_sample=1, loss_function='mae', estimator='median')
clf1.fit(x_krr_train, y_krr_train)

clf2 = Regression_Tree(max_depth=2, min_sample=1, loss_function='mae', estimator='median')
clf2.fit(x_krr_train, y_krr_train)

clf3 = Regression_Tree(max_depth=3, min_sample=1, loss_function='mae', estimator='median')
clf3.fit(x_krr_train, y_krr_train)

clf4 = Regression_Tree(max_depth=4, min_sample=1, loss_function='mae', estimator='median')
clf4.fit(x_krr_train, y_krr_train)

clf5 = Regression_Tree(max_depth=5, min_sample=1, loss_function='mae', estimator='median')
clf5.fit(x_krr_train, y_krr_train)

clf6 = Regression_Tree(max_depth=6, min_sample=1, loss_function='mae', estimator='median')
clf6.fit(x_krr_train, y_krr_train)

plot_size = 0.001
x_range = np.arange(0., 1., plot_size).reshape(-1, 1)

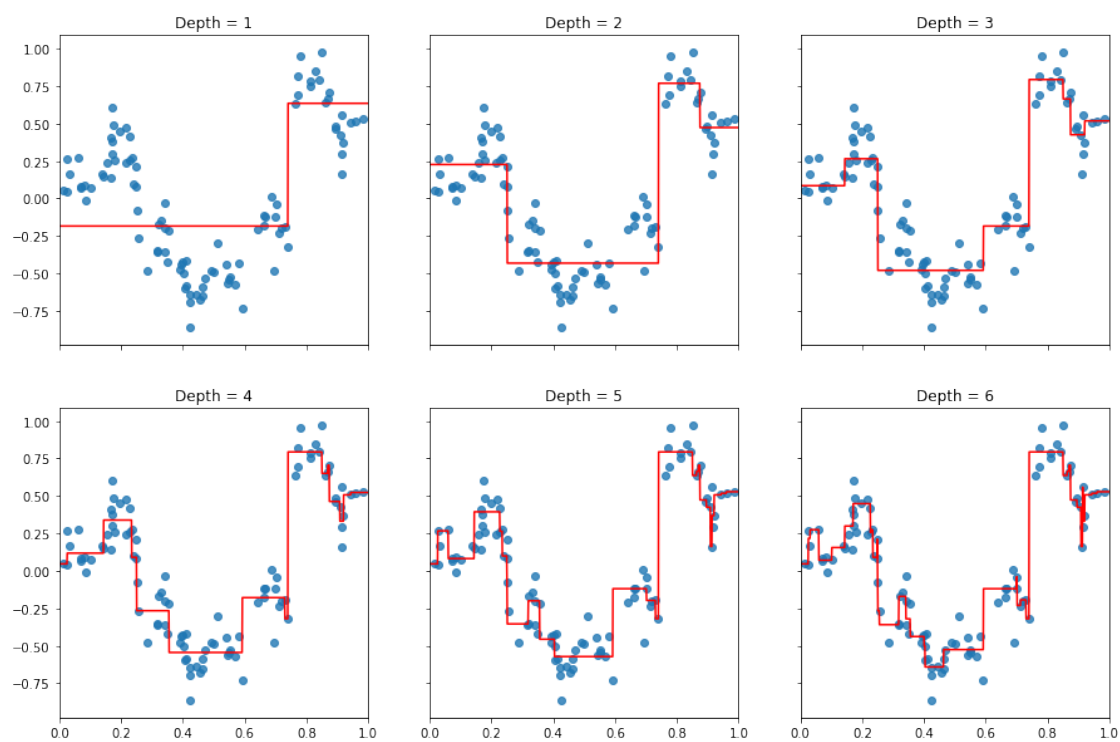
f2, axarr2 = plt.subplots(2, 3, sharex='col', sharey='row', figsize=(15, 10))

for idx, clf, tt in zip(product([0, 1], [0, 1, 2]),
                        [clf1, clf2, clf3, clf4, clf5, clf6],
                        ['Depth = {}'.format(n) for n in range(1, 7)]):

    y_range_predict = np.array([clf.predict_instance(x) for x in x_range]).reshape(-1, 1)

    axarr2[idx[0], idx[1]].plot(x_range, y_range_predict, color='r')
    axarr2[idx[0], idx[1]].scatter(x_krr_train, y_krr_train, alpha=0.8)
    axarr2[idx[0], idx[1]].set_title(tt)
    axarr2[idx[0], idx[1]].set_xlim(0, 1)
plt.show()

/home/cfizette/anaconda3/envs/ml/lib/python3.6/site-packages/numpy/core/fromnumeric.py:3118: RuntimeWarning:
    out=out, **kwargs)
/home/cfizette/anaconda3/envs/ml/lib/python3.6/site-packages/numpy/core/_methods.py:85: RuntimeWarning:
    ret = ret.dtype.type(ret / rcount)
/home/cfizette/anaconda3/envs/ml/lib/python3.6/site-packages/ipykernel/__main__.py:11: RuntimeWarning:
```



8 Gradient Boosting Method

```
In [15]: #Pseudo-residual function.  
         #Here you can assume that we are using L2 loss  
  
         def pseudo_residual_L2(train_target, train_predict):  
             '''  
             Compute the pseudo-residual based on current predicted value.  
             '''  
             return train_target - train_predict  
  
         class ConstantModel(BaseEstimator, RegressorMixin):  
             def __init__(self, c):  
                 self.c = c  
  
             def fit(self, x, y=None):  
                 pass  
  
             def predict(self, x):  
                 return self.c * np.ones(len(x))  
  
In [18]: class gradient_boosting():  
         '''  
         Gradient Boosting regressor class  
         :method fit: fitting model  
         '''  
  
         def __init__(self, n_estimator, pseudo_residual_func, learning_rate=0.1, min_sample  
             '''  
             Initialize gradient boosting class  
  
             :param n_estimator: number of estimators (i.e. number of rounds of gradient boo  
             :pseudo_residual_func: function used for computing pseudo-residual  
             :param learning_rate: step size of gradient descent  
             '''  
  
             self.n_estimator = n_estimator  
             self.pseudo_residual_func = pseudo_residual_func  
             self.learning_rate = learning_rate  
             self.min_sample = min_sample  
             self.max_depth = max_depth  
             self.estimators = [ConstantModel(c=0)]  
  
         def calc_pseudo_residual(self, X, y):  
             return self.predict(X) - y.flatten()  
  
         def fit(self, train_data, train_target):  
             '''  
             Fit gradient boosting model  
             '''
```

```

for _ in range(self.n_estimator):
    pseudo_residual = self.calc_pseudo_residual(train_data, train_target)
    #pdb.set_trace()
    new_estimator = DecisionTreeRegressor(max_depth=self.max_depth,
                                          min_samples_split=self.min_sample)
    new_estimator.fit(X=train_data, y=-pseudo_residual)
    #pdb.set_trace()
    self.estimators.append(new_estimator)

def predict(self, test_data):
    '''
    Predict value
    '''
    preds = np.zeros(len(test_data))
    for estimator in self.estimators:
        preds += self.learning_rate * estimator.predict(test_data)
    return preds

```

9 2-D GBM visualization - SVM data

```
In [19]: # Plotting decision regions
x_min, x_max = x_train[:, 0].min() - 1, x_train[:, 0].max() + 1
y_min, y_max = x_train[:, 1].min() - 1, x_train[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))


f, axarr = plt.subplots(2, 3, sharex='col', sharey='row', figsize=(10, 8))

for idx, i, tt in zip(product([0, 1], [0, 1, 2]),
                     [1, 5, 10, 20, 50, 100],
                     ['n_estimator = {}'.format(n) for n in [1, 5, 10, 20, 50, 100]]):

    gbt = gradient_boosting(n_estimator=i, pseudo_residual_func=pseudo_residual_L2, max
    gbt.fit(x_train, y_train)

    Z = np.sign(gbt.predict(np.c_[xx.ravel(), yy.ravel()]))
    Z = Z.reshape(xx.shape)

    axarr[idx[0], idx[1]].contourf(xx, yy, Z, alpha=0.4)
    axarr[idx[0], idx[1]].scatter(x_train[:, 0], x_train[:, 1], c=y_train_label.flatten
    axarr[idx[0], idx[1]].set_title(tt)
```

10 1-D GBM visualization - KRR data

```
In [22]: plot_size = 0.001
         x_range = np.arange(0., 1., plot_size).reshape(-1, 1)


         f2, axarr2 = plt.subplots(2, 3, sharex='col', sharey='row', figsize=(15, 10))

         for idx, i, tt in zip(product([0, 1], [0, 1, 2]),
                                [1, 5, 10, 20, 50, 100],
                                ['n_estimator = {}'.format(n) for n in [1, 5, 10, 20, 50, 100]]):

             gbm_1d = gradient_boosting(n_estimator=i, pseudo_residual_func=pseudo_residual_L2,
                                         gbm_1d.fit(x_krr_train, y_krr_train)

             y_range_predict = gbm_1d.predict(x_range)

             axarr2[idx[0], idx[1]].plot(x_range, y_range_predict, color='r')
             axarr2[idx[0], idx[1]].scatter(x_krr_train, y_krr_train, alpha=0.8)
             axarr2[idx[0], idx[1]].set_title(tt)
             axarr2[idx[0], idx[1]].set_xlim(0, 1)
```



multiclass-skeleton-code

April 29, 2019

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets.samples_generator import make_blobs
import pdb
from collections.abc import Iterable
from tqdm import tqdm
```

```
%matplotlib inline
```


```
In [2]: # Create the training data
```

```
np.random.seed(2)
```

```
X, y = make_blobs(n_samples=300, cluster_std=.25, centers=np.array([(-3,1),(0,2),(3,1)]))
```

```
plt.scatter(X[:, 0], X[:, 1], c=y, s=50)
```

```
Out[2]: <matplotlib.collections.PathCollection at 0x7fa982119470>
```



```

In [3]: from sklearn.base import BaseEstimator, ClassifierMixin, clone

class OneVsAllClassifier(BaseEstimator, ClassifierMixin):
    """
    One-vs-all classifier
    We assume that the classes will be the integers 0,...,(n_classes-1).
    We assume that the estimator provided to the class, after fitting, has a "decision_function"
    returns the score for the positive class.
    """
    def __init__(self, estimator, n_classes):
        """
        Constructed with the number of classes and an estimator (e.g. an
        SVM estimator from sklearn)
        @param estimator : binary base classifier used
        @param n_classes : number of classes
        """
        self.n_classes = n_classes
        self.estimators = [clone(estimator) for _ in range(n_classes)]
        self.fitted = False

    def fit(self, X, y=None):
        """
        This should fit one classifier for each class.
        self.estimators[i] should be fit on class i vs rest
        @param X: array-like, shape = [n_samples,n_features], input data
        @param y: array-like, shape = [n_samples,] class labels
        @return returns self
        """
        for class_n, estimator in enumerate(self.estimators):
            labels = (y == class_n).astype(int)
            estimator.fit(X, labels)
        self.fitted = True
        return self

    def decision_function(self, X):
        """
        Returns the score of each input for each class. Assumes
        that the given estimator also implements the decision_function method (which sklearn
        and that fit has been called.
        @param X : array-like, shape = [n_samples, n_features] input data
        @return array-like, shape = [n_samples, n_classes]
        """
        if not self.fitted:
            raise RuntimeError("You must train classifier before predicting data.")

        if not hasattr(self.estimators[0], "decision_function"):
            raise AttributeError(
                "Base estimator doesn't have a decision_function attribute.")

```

```

        return np.array([e.decision_function(X) for e in self.estimators]).T

    def predict(self, X):
        """
        Predict the class with the highest score.
        @param X: array-like, shape = [n_samples,n_features] input data
        @returns array-like, shape = [n_samples,] the predicted classes for each input
        """
        return np.argmax(self.decision_function(X), axis=1)

```

```

In [4]: #Here we test the OneVsAllClassifier
from sklearn import svm
svm_estimator = svm.LinearSVC(loss='hinge', fit_intercept=False, C=200)
clf_onevsall = OneVsAllClassifier(svm_estimator, n_classes=3)
clf_onevsall.fit(X,y)

for i in range(3) :
    print("Coeffs %d"%i)
    print(clf_onevsall.estimators[i].coef_) #Will fail if you haven't implemented fit ye

# create a mesh to plot in
h = .02 # step size in the mesh
x_min, x_max = min(X[:,0])-3,max(X[:,0])+3
y_min, y_max = min(X[:,1])-3,max(X[:,1])+3
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))
mesh_input = np.c_[xx.ravel(), yy.ravel()]

Z = clf_onevsall.predict(mesh_input)
Z = Z.reshape(xx.shape)
plt.contourf(xx, yy, Z, cmap=plt.cm.coolwarm, alpha=0.8)
# Plot also the training points
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.coolwarm)

from sklearn import metrics
metrics.confusion_matrix(y, clf_onevsall.predict(X))

```


```

Coeffs 0
[[-1.05852964 -0.90293998]]
Coeffs 1
[[ 0.31047003 -0.19010007]]
Coeffs 2
[[ 0.8916347  -0.82599479]]

```

```
/home/cfizette/anaconda3/envs/ml/lib/python3.6/site-packages/sklearn/svm/base.py:931: ConvergenceWarning:
  "the number of iterations.", ConvergenceWarning)
```

```
Out[4]: array([[100,  0,  0],
               [ 0, 100,  0],
               [ 0,  0, 100]])
```



Multiclass SVM

```
In [16]: def zeroOne(y,a) :
        '''
        Computes the zero-one loss.
        @param y: output class
        @param a: predicted class
        @return 1 if different, 0 if same
        '''
        return int(y != a)

def featureMap(X,y,num_classes) :
    '''
    Computes the class-sensitive features.
    @param X: array-like, shape = [n_samples,n_inFeatures] or [n_inFeatures,], input fe
    @param y: a target class (in range 0,..,num_classes-1)
    @return array-like, shape = [n_samples,n_outFeatures], the class sensitive features
    '''
    #The following line handles X being a 1d-array or a 2d-array
    num_samples, num_inFeatures = (1,X.shape[0]) if len(X.shape) == 1 else (X.shape[0],
    if num_samples == 1:
        X = [X]

    # n_out_features = n_classes * n_in_features
    # feature_map[y*n_in_features : y*n_in_features + n_in_features] = X

    feature_map = np.zeros((num_samples, num_inFeatures * num_classes))

    for i, xi in enumerate(X):
        start = y*num_inFeatures
        end = start + num_inFeatures
        feature_map[i,start:end] = xi

    '''
    for i in range(num_samples):
        xi = X[i]
        yi = y
        start = yi*num_inFeatures
        end = start + num_inFeatures
        feature_map[i,start:end] = xi
    '''
    return feature_map

def sgd(X, y, num_outFeatures, subgd, eta = 0.1, T = 10000):
    '''
    Runs subgradient descent, and outputs resulting parameter vector.
    @param X: array-like, shape = [n_samples,n_features], input training data
    @param y: array-like, shape = [n_samples,], class labels
```

```

    @param num_outFeatures: number of class-sensitive features
    @param subgd: function taking x,y and giving subgradient of objective
    @param eta: learning rate for SGD
    @param T: maximum number of iterations
    @return: vector of weights
    '''
    num_samples = X.shape[0]
    w = np.zeros(num_outFeatures)
    for _ in tqdm(range(T)):
        for xi, yi in zip(X, y):
            w -= eta * subgd(xi, yi, w)
    return w

class MulticlassSVM(BaseEstimator, ClassifierMixin):
    '''
    Implements a Multiclass SVM estimator.
    '''
    def __init__(self, num_outFeatures, lam=1.0, num_classes=3, Delta=zeroOne, Psi=feat
    '''
    Creates a MulticlassSVM estimator.
    @param num_outFeatures: number of class-sensitive features produced by Psi
    @param lam: l2 regularization parameter
    @param num_classes: number of classes (assumed numbered 0,..,num_classes-1)
    @param Delta: class-sensitive loss function taking two arguments (i.e., target
    @param Psi: class-sensitive feature map taking two arguments
    '''
    self.num_outFeatures = num_outFeatures
    self.lam = lam
    self.num_classes = num_classes
    self.Delta = Delta
    self.Psi = lambda X,y : Psi(X,y,num_classes)
    self.fitted = False

    def generalized_hinge_loss(self, x, y, w):
        return max([self.class_hinge_loss(x,y,w,y_pred) for y_pred in range(self.n_clas

    def class_hinge_loss(self, x, y, w, y_pred):
        return self.Delta(y, y_pred) + (self.Psi(x, y_pred) - self.Psi(x, y))@w

    def subgradient(self,x,y,w):
        '''
        Computes the subgradient at a given data point x,y
        @param x: sample input
        @param y: sample class
        @param w: parameter vector
        @return returns subgradient vector at given x,y,w
        '''
        y_hat = np.argmax([self.class_hinge_loss(x,y,w,y_pred) for y_pred in range(self

```



```

sgd = 2 * self.lam * w + self.Psi(x,y_hat).flatten() - self.Psi(x,y).flatten()
return sgd

def fit(self,X,y,eta=0.1,T=10000):
    """
    Fits multiclass SVM
    @param X: array-like, shape = [num_samples,num_inFeatures], input data
    @param y: array-like, shape = [num_samples,], input classes
    @param eta: learning rate for SGD
    @param T: maximum number of iterations
    @return returns self
    """
    self.coef_ = sgd(X,y,self.num_outFeatures,self.subgradient,eta,T)
    self.fitted = True
    return self

def decision_function(self, X):
    """
    Returns the score on each input for each class. Assumes
    that fit has been called.
    @param X : array-like, shape = [n_samples, n_inFeatures]
    @return array-like, shape = [n_samples, n_classes] giving scores for each sample
    """
    if not self.fitted:
        raise RuntimeError("You must train classifier before predicting data.")

    return np.array([self.Psi(X, y)@self.coef_ for y in range(self.num_classes)]).T

def predict(self, X):
    """
    Predict the class with the highest score.
    @param X: array-like, shape = [n_samples, n_inFeatures], input data to predict
    @return array-like, shape = [n_samples,], class labels predicted for each data
    """
    raw_pred = self.decision_function(X)
    pred = np.argmax(raw_pred, axis=1)
    return pred

```

```

In [17]: #the following code tests the MulticlassSVM and sgd
         #will fail if MulticlassSVM is not implemented yet
         est = MulticlassSVM(6,lam=1)
         est.fit(X,y, T=500, eta=0.01)
         print("w:")
         print(est.coef_)
         Z = est.predict(mesh_input)
         Z = Z.reshape(xx.shape)
         plt.contourf(xx, yy, Z, cmap=plt.cm.coolwarm, alpha=0.8)
         # Plot also the training points

```

```
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.coolwarm)
```

```
from sklearn import metrics  
metrics.confusion_matrix(y, est.predict(X))
```

```
100%|| 500/500 [00:10<00:00, 46.90it/s]
```

w:

```
[-0.34533686 -0.04675874  0.02345881  0.06871364  0.32187805 -0.02195491]
```

```
Out[17]: array([[100,  0,  0],  
               [ 0, 100,  0],  
               [ 0,  0, 100]])
```

